Supplementary Material for "Class-agnostic candidate generation and calendar-derived event windows for rapid landslide mapping"

S1-A. Implementation details

Encoders. Optical and SAR branches use Swin-B with out_indices (1, 2, 3) producing 256/512/1024 channels for optical and 192/384/768 for SAR. Topography (5 channels) and change (3 or 6 channels) use the lightweight modality stem: two stride-2 convolutions followed by three residual stages, producing the same three level widths. The fused representation is 256/512/1024 after per-level 1x1 projection.

Fusion. Per level, the projected modality features are stacked along the modality axis and passed through a 1x1 bottleneck (GELU) to per-modality logits, which are masked by the per-sample modality mask and softmaxed along the modality axis. The fused level is the gate-weighted sum. Gate logits are zero initialised.

Training. Eight epochs, AdamW, learning rate 1e-5, weight decay 1e-4, gradient clipping at norm 0.1, automatic mixed precision, batch size 4, resolution 320 (native patches are 128). Parameter-wise multipliers: backbone 0.5, modality encoders 1.0, fusion 1.0, head 1.0, router 1.0, meta-net 0.1, decoder 0.0, learnable text embedding 0.1. Only the most recent checkpoint is kept.

Verifier. MLP with hidden widths 256 and 64, listwise softmax objective over candidates of the same image, input 523 dimensions (frozen CLIP ViT-B/32 embedding plus 11 within-box statistics), Adam at 1e-3 with a step scheduler (step 10, gamma 0.7). Model-native scores are deliberately excluded from the input.

Interval intervention (Section 4.3 of the main text). Training stacks are built as delta NDVI, delta NDWI and delta NBR between the bracketing acquisitions, widened by one acquisition on each side. The deployment stack is unchanged from the long baseline. Only the training manifest's change paths change; the processor normalisation is recomputed from the training split.

S1-B. Full ablation (A0 to A9)

Table S1-B. Architecture ablation, AR@900. Reference values are from the 2026-10-12 batch; two configurations were re-run on 2026-10-20 with the published recipe (batch 4, eight epochs, seed 42) and are marked "re-run". In the last column, not re-run marks configurations that were not re-trained for this revision; only A0, A4 and A9 were, and the re-run values are reported here next to the archived reference.

Channel-count control. The control arm keeps the architecture, the optimiser and the candidate budget of the widened-window arm and changes only how the three additional change channels are filled: the three long-baseline channels (dNDVI, dNDWI, dNBR) are repeated once each, so the arm has the same six change channels as the fused arm but carries no event-window information. It reaches AR@900 0.7041 (AR@300 0.5857, AR@100 0.4384), against 0.7094 for the anchor, so the extra channels by themselves neither add nor remove recall; the gain of the fused arm (0.7583 over two seeds) therefore comes from the interval those channels carry.

Arm	Modalities	Reference AR@900	Re-run AR@900
A0	optical	0.3533	**0.3714**
A1	optical + topography	0.4220	not re-run
A2	optical + SAR	0.3745	not re-run
A3	optical + SAR + topography	0.4574	not re-run
A4	A3 + change	0.7167	0.7094 (anchor)
A5	A4 + modality tokens	0.6905	not re-run
A6	A5 + cross-modal router	0.7059	not re-run
A7	A6 without router	0.7029	not re-run
A8	A6 without SAR	0.6729	not re-run
A9	A6 without change	0.4799	**0.4337**

Reading A0. The re-run value lies within the run-to-run band of the reference (+0.018 against a band of 0.025), and its 95% CI, [0.350, 0.392], contains the reference value. A0 also reproduces the deployment-side picture: without change channels the strict-budget recall collapses to 0.092.

Reading A9. The re-run value is **below** the reference band (-0.046). The configuration was verified to apply the intended overrides (SAR and topography on, change off, modality tokens and router on) and to complete all eight epochs. A9 contains the two components that the paper reports as not improving performance, so its run-to-run spread is expected to be larger than A4's. We report the discrepancy rather than adjusting the arm, and treat the 2026-10-12 value as the archived reference. This is an open reproducibility gap for that single configuration.

S1-C. Candidate budgets and cost

Table S1-C1. Per-arm recall on the four-event split, k = 100/300/900 proposals per image. AR@k is image-averaged (the convention used for every ratio in the main text); micro is ground-truth weighted. Both conventions are reported for every arm so that no ratio in the paper mixes them.

Arm	AR@k, image-averaged			micro recall@k (ground-truth weighted)			Realised budget
	100	300	900	100	300	900	
PF-RPN (prompt-free)	0.4652	0.5908	0.7059	0.3170	0.4907	0.6619	900
Fine-tuned YOLO-World	0.0606	0.1399	0.3321	0.0431	0.1131	0.2900	900
Closed-set YOLOv8s	0.0617	0.1364	0.3203	0.0405	0.1046	0.2768	900
Grounding DINO, zero-shot	0.0462	0.0993	0.1828	0.0341	0.0769	0.1595	831-900
YOLO-World, zero-shot	0.0121	0.0315	0.1019	0.0076	0.0248	0.0762	816-869

Notes: the four baseline arms were scored through the shared 900-candidate export pipeline; PF-RPN values are from the in-memory evaluator on the same 1,759 images / 4,227 ground-truth boxes (the route fills the 900 cap, so the two paths coincide at k = 900). Grounding DINO is the fair export (baselines/fair/gdino_test, AR@900 = 0.182789); a second export of the same checkpoint (gdino_main_fair, 0.182603) differs only beyond rank 100 and is deprecated. Realised budgets after the minimum-box filter: PF-RPN 900, Grounding DINO about 831, YOLO-World zero-shot about 869; the two fine-tuned arms fill the 900 cap.

S1-C2. False positives per image normalisation (10 FPPI). PF-RPN 0.1270; Grounding DINO 0.0019; fine-tuned YOLO-World 0.0021. Ratios 34.5x and 41.4x. Earlier drafts quoted 67x and 60x from a superseded normalisation; those values are retired and do not appear in the main text.

S1-D. Reproducibility of the criterion

Table S1-D. Per-fold criterion values and measured recall, sorted by separability. The criterion is the label-free separability AUC; the measured recall is PF-RPN AR@900 on that fold; the gap is PF-RPN over Grounding DINO zero-shot at the same 900-candidate budget.

Held-out fold	criterion AUC	measured AR@900	AR@900 gap (PF/GDINO)
Kyrgyzstan1 (LOEO)	0.4635	0.7592	1.57
Central Asia (LORO)	0.4771	0.6831	1.54
New Zealand (LOEO)	0.4922	0.6911	1.37
East Asia (LORO)	0.7052	0.7319	3.78
Indonesia (LOEO)	0.7054	0.7839	4.68
Hiroshima (LOEO)	0.7059	0.7366	4.97
Hokkaido (LOEO)	0.7317	0.7350	3.71
South-East Asia (LORO)	0.7397	0.7697	4.29
Italy (LOEO)	0.8323	0.2906	5.65

Chimanimani (LOEO)	0.9064	0.7369	4.75

Summary statistics: Pearson 0.915, Spearman 0.758, 95% CI [0.682, 0.986], leave-one-fold-out 0.856, sign agreement 10 of 10, largest single-fold influence 0.041, correlation with absolute recall -0.266 (CI crosses zero). Two-tier behaviour: AUC at or above 0.70 predicts a gap of 3.71 to 5.65 (17.7% to 26.9% of the detector's recall recovered), AUC at or below 0.50 predicts 1.37 to 1.57 (63.7% to 72.8%); no fold falls between 0.492 and 0.705.

S1-E. Change-channel interventions (all arms)

Table S1-E1. Same-tile permutation, anomaly-free controls (cloud, anchor 0.7059).

Arm	Operation	AR@900	Difference
clean	none	0.7059	anchor
change permuted	spatial permutation within tile	0.3866	-0.319
SAR permuted	spatial permutation within tile	0.6911	-0.015
topography permuted	spatial permutation within tile	0.6991	-0.007

Table S1-E2. Temporal interventions (cloud, anchor 0.7059).

Arm	Window	Span	AR@900	Difference
clean	event, first to last	1.00 T	0.7059	anchor
temporal mispairing	two earliest acquisitions	about 0.02 T	0.3535	-0.352
equal-span control	outside window, maximum available	0.48 T	0.4632	-0.243

Strict equal-span control is not constructible on this dataset: the true pre/post span is the full series, and the annotated window lies in its middle. This is why the interval question is resolved by the calendar analysis of Section 3.2 of the main text rather than by span matching.

Table S1-E3. Inference-time substitution on a frozen model (RTX 4060, anchor 0.7085).

Change stack given to the model	Normalisation	AR@100	AR@900	Difference at 900
long baseline (matched)	model's own	0.4655	0.7085	anchor
event pair	model's own	0.2711	0.5460	-0.163
widened event window	model's own	0.3712	0.6217	-0.087
widened event window	variant statistics	0.3114	0.5941	-0.114
event pair	variant statistics	0.2362	0.5119	-0.197

S1-E4. (exploratory arm definition) Two intervals in one input (exploratory). A six-channel change stack that concatenates the long baseline and the widened event window was built for every tile, with the existing fusion gate left untouched, so that the three extra channels can be varied independently of the interval. Repeating the deployed long-baseline stack in those channels, and training and deploying on the long baseline, leaves recall on the anchor: AR@900 0.7041 against 0.7094 on the same machine (95% CI [0.6822, 0.7246]). Placing the event window in the same three channels, so that both intervals enter the model, reaches 0.7583 over two seeds (0.7634 and 0.7532, s.d. 0.0072), whose confidence interval does not overlap the control's. Against the widened-event training arm of the same seed, a paired per-image bootstrap gives +0.049 at k=100 (95% CI [+0.025, +0.073]) and +0.029 at k=300 ([+0.008, +0.050]), but +0.007 at k=900 ([-0.011, +0.026]). Unlike the widened-event arm, this configuration needs the event stream at inference, so it is reported as an exploratory extension rather than as part of the free intervention (Table S1-E5).

Table S1-E5. Six-channel arms (exploratory). The control repeats the deployed long-baseline stack in the three extra change channels; the two-interval arm places the widened event window there. All arms are scored with 900 candidates per image on the four-event test split (1,759 images) on the RTX 3080 Ti,

whose own anchor is 0.7094; the archived anchor of 0.7059 sits 0.005 lower, inside the run-to-run noise of 0.0059.

Arm	AR@100	AR@300	AR@900
six channels, both intervals long (control)	0.4384	0.5857	0.7041
six channels, long + event window, seed 42	0.5789	0.6867	0.7634
six channels, long + event window, seed 123	0.5607	0.6783	0.7532
mean and sd (two seeds)	0.5698 ± 0.0129	0.6825 ± 0.0059	0.7583 ± 0.0072

S1-F. Interval analysis (pixel level and calendar)

Table S1-F1. Pixel-level separability by window. Four test events, about 400 annotated tiles, official masks as ground truth, paired bootstrap over tiles.

Window	**Span**	**AUC**	**recall@1%FPR**	**usable pixels**	**delta AUC vs long baseline**
long baseline	360 d	0.7771	0.1948	0.982	0.000 (reference)
event pair	10 d	0.7439	0.1189	0.776	-0.045 [-0.071, -0.055]
event, widened by one	25 d	**0.8476**	**0.2704**	**0.985**	**+0.066 [+0.046, +0.060]**
event, widened by two	70 d	0.8315	0.2742	0.987	+0.051 [+0.031, +0.044]
widened, validity-gated	70 d	0.8303	0.2709	0.990	+0.050 [+0.028, +0.044]
same span, displaced	25 d	0.5064	0.0155	0.960	-0.240 [-0.273, -0.256]
label-free NDVI-drop rule	10 d	0.5147	0.0300	0.875	-0.259 [-0.290, -0.270]

Table S1-F2. Calendar reproduction of the annotated window, per inventory.

Inventory	**n**	**exact**	**within one acquisition**
italy	1,799	1.000	1.000
chimanimani	345	0.986	1.000
hiroshima	324	1.000	1.000
kyrgyzstan1	186	0.823	1.000
hokkaido	144	1.000	1.000
newzealand	94	0.926	1.000
itogon	82	1.000	1.000
kyrgyzstan2	58	0.759	1.000
indonesia	47	0.979	1.000
china	28	0.821	1.000
usa	28	1.000	1.000
all	**3,135**	**0.979**	**1.000**

Table S1-F3. Distance from the event date to the bracketing acquisitions, over the annotated tiles. The annotated window is a single adjacent acquisition pair in 100% of tiles, positioned at a median relative index of 0.50.

Bracketing acquisition	**median (days)**	**90th percentile (days)**
previous	3	5
next	4	7

S1-F4. Window properties and the gain, by protocol

Table S1-F4. Window properties and the interval gain, by protocol. Medians over the annotated tiles of each split; "tiles with dates" counts the tiles whose annotation carries a pre/post pair, and only those contribute to the window columns. The last two rows score the same 253 Indonesian tiles and differ only in the backing model, so the sign of the gain is set by the protocol rather than by the region.

Split / event	GT tiles	cadence (d)	event ±1 span (d)	usable long / pair / ±1	anchor AR@900	±1 AR@900	Δ AR@900
four-event test, chimanimani	345	22	25	0.999 / 0.851 / 1.000	0.7091	0.7706	+0.0615
four-event test, china	47	101	175	0.926 / 0.000 / 0.979	0.7963	0.8102	+0.0138
four-event test, indonesia	253	88	170	1.000 / 0.179 / 1.000	0.7137	0.7551	+0.0414
four-event test, usa	28	28	50	0.991 / 0.822 / 0.995	0.5284	0.4991	-0.0294
leave-one-region-out fold, indonesia	253	88	170	1.000 / 0.179 / 1.000	0.7891	0.7717	-0.0173

The last two rows of Table S1-F4 hold the region, its acquisition calendar and its windows fixed and change only the backing model: under the event-level split the same 253 Indonesian tiles gain 0.0414 at AR@900, under the leave-one-region-out fold they lose 0.0173. Window validity is not the difference, because the widened window is fully usable in both (1.000), and neither is cadence (88 d in both, an event window of 170 d). What differs is how strong the baseline already is on that region, 0.7137 against 0.7891, which is consistent with the gain being concentrated where candidates are scarce. The intervention therefore does not replicate under the second protocol, and Section 5.3 of the main text of the paper reports it as an open limitation rather than as a boundary that has been explained away. Note also that the annotated pair is almost entirely unusable in Indonesia and china (0.179 and 0.000), which is the same effect that makes the single adjacent pair the weakest window in Table S1-F1.

S1-G. Annotation-date quality

Table S1-G. Multi-event tiles among the 4,392 positive tiles.

Candidate dates	n	share	median GT pixels	median date confidence
single window	3,139	71.5%	224	1.000
2	366	8.3%	245	0.444
3	273	6.2%	385	0.464
4	177	4.0%	446	0.451
5	109	2.5%	651	0.498
6	83	1.9%	753	0.488
7	96	2.2%	1,110	0.526
8	61	1.4%	1,255	0.579
9	33	0.8%	1,907	0.575
10	27	0.6%	2,018	0.599
11 or more	28	0.6%	2,510 (at 11)	0.610

Rule used for multi-event tiles: the highest-confidence candidate (first entry of the ranked list). Unannotated tiles: the neighbouring event's calendar, or the series midpoint when no event date exists.

S1-H. Reproduction and validation

Table S1-H1. Data and pipeline checks.

Check	Result
Tile and array count	10,015 tiles, 60,090 arrays, 0 inconsistencies
Long-baseline reproduction	40 tiles across 4 events, maximum absolute difference 0, masks identical 40 of 40
Anchors by machine	0.7059 (cloud), 0.7085 (RTX 4060), 0.7094 (RTX 3080 Ti)
Run-to-run envelope per fold	0.0076, 0.0234, 0.0251
QA gate	0 flagged artefacts

Table S1-H2. Claims added in this revision and their artefacts.

Claim	Value	Artefact
Calendar reproduction	0.979 / 1.000, n=3,135	calendar rule output
Pixel-level window AUC	0.7771 / 0.7439 / 0.8476 / 0.5064	window mechanics table
Frozen substitution	0.7085 to 0.6217 / 0.5460	substitution evaluations
Training-axis ordering	0.6969 / 0.7094 / 0.7563	Table Z
Paired gains at three budgets	+0.0644 / +0.0694 / +0.0468	per-image dumps, paired bootstrap
Multi-event share	0.285	feature dump
Channel-count control	0.7041 against 0.7094 on the same machine	6-channel duplicated-stack run
Two-interval representation	0.7583 ± 0.0072, two seeds	6-channel arm, seeds 42 and 123
Cross-protocol fold (negative)	0.7891 to 0.7717 at AR@900; -0.0423 at AR@300	leave-one-region-out run, paired bootstrap

S1-I. Re-run values for the ablation table

Table S1-I. Re-runs performed for this revision. All carry batch size 4, eight epochs, seed 42, and the published recipe.

Configuration	Reference	Re-run	Difference	Status
A0	0.3533	0.3714	+0.018	within band, CI contains reference
A9	0.4799	0.4337	-0.046	**outside band, open discrepancy**
A4 (anchor)	0.7059	0.7094	+0.0035	within band

Seed replicates of the widened-event training arm. Seeds 42, 123 and 777 were run on the identical data definition and are evaluated on the deployed long-baseline root, so only the training interval differs from the anchor. The earlier seed-42 checkpoint was trained before the multi-event window rule was finalised, which is why seed 42 is repeated here.

Seed	AR@100	AR@300	AR@900
42	0.5165	0.6456	0.7594
123	0.5127	0.6459	0.7400
777	0.5299	0.6551	0.7429
mean and sd	0.5197 ± 0.0091	0.6488 ± 0.0054	0.7474 ± 0.0105

The three seeds give AR@900 0.7594, 0.7400 and 0.7429, a mean of 0.7474 with a seed standard deviation of 0.0105; at k=300 the mean is 0.6488 (s.d. 0.0054) and at k=100 0.5197 (s.d. 0.0091). Combining the

target's seed variance with the anchor's run-to-run variance of 0.0059 puts the AR@900 gain at 5.5 standard errors, the AR@300 gain at 10.6 and the AR@100 gain at 8.6; a paired per-image bootstrap on the first seed gives +0.0468 (95% CI [+0.029, +0.065]) at k=900.

S1-J. Pre-registered test on the weakest fold

Why this test. The leave-one-region-out fold on Indonesia loses 0.0173 AR@900 under the widened-event window, whereas the event-level split gains 0.0414 on the same 253 Indonesian tiles; the two protocols hold the region, its calendar and its windows fixed and differ in the backing model, whose anchor is 0.7137 against 0.7891 on that subset (Table S1-F4). One reading is that the intervention pays in proportion to the headroom a model leaves on the deployment region. This section fixes the test of that reading, and its decision rule, before the runs finish.

Prior evidence that does not favour the account. The same four-event split gives no such gradient: across its four events the anchor level does not predict the gain (Pearson +0.65 on the values and +0.20 on ranks), the event with the lowest anchor (usa, 0.528) is the one that loses most (-0.0294 at AR@900), and the event with the highest anchor (china, 0.796) gains only +0.0138. The account is therefore tested as a candidate explanation that the existing evidence does not support, which is why a positive outcome would be reported as surprising rather than as confirmation.

Design. Italy is the fold with the weakest archived anchor (AR@900 0.2896, re-scored here at 0.2950, a difference inside the run-to-run envelope). The widened-event arm is trained for eight epochs on the fold's identical 5,716-tile training set with seeds 42, 123 and 777, and evaluated on the fold's long-baseline root; the archived anchor weights are re-scored on the same root. Preprocessing, the recomputation of the change-channel normalisation from the training split, and the evaluation protocol follow Section 4.3 of the main text, so the only difference from the anchor is the change stack used in training. No leave-one-out fold has an anchor between 0.5 and 0.6 - the archived anchors are bimodal, 0.2896 and then 0.6757 to 0.7889 - so a multi-point dose-response gradient cannot be assembled from the archived folds and a single fold is what this design can supply.

Decision rule, fixed in advance. Let dAR@k be the seed-averaged per-image AR@k minus the anchor's per-image AR@k, and let the decision interval be the 95% percentile interval of a paired bootstrap over the test images that carry ground truth (20,000 resamples). At k = 900: (a) if the upper bound is at or below 0, the headroom account is dropped from the paper and the result is reported here as a negative control; (b) if the lower bound is above 0, the account is reported as consistent with the data and not as confirmed, because one fold cannot separate headroom from event difficulty, annotation quality or domain shift; (c) if the interval contains 0, the account is downgraded to a speculation and the paper states only the conditional form.

Reporting obligation. Per-seed intervals are reported alongside the seed-averaged interval, and the seed dispersion is combined with the anchor's run-to-run standard deviation of 0.0059 in the same way as Section 4.3 of the main text. Whatever the outcome, the test is reported: a null or negative result enters this document and one sentence of Section 5.3 of the main text rather than being withheld. What the rule freezes is the claim, not the data.

S1-K. Result of the pre-registered test (negative control)

The decision rule fixed in S1-J was applied to the three budgets in turn. The seed-averaged intervention lies below the anchor at every budget and the upper end of every 95% interval lies below zero, so the pre-registered termination clause applies as written: the headroom account is withdrawn rather than re-specified after the fact.

Table S1-K. Pre-registered test on the Italy leave-one-event-out fold. The anchor is the archived long-baseline model, re-scored on this machine (1,799 ground-truth boxes); the arm is the same architecture trained on the widened event window and deployed on the long baseline; intervals are paired per-image bootstrap intervals over the fold's images.

k	anchor AR@k	seed-averaged AR@k	delta	95% CI	verdict
100	0.1392	0.0944	-0.0448	[-0.0498, -0.0400]	KILL
300	0.2004	0.1484	-0.0520	[-0.0576, -0.0465]	KILL
900	0.2950	0.2280	-0.0670	[-0.0734, -0.0607]	KILL

Per-seed values behind the seed-averaged row: AR@900 of 0.2303, 0.2367 and 0.2172 against the anchor's 0.2950, that is deltas of -0.0648, -0.0584 and -0.0779 for seeds 42, 123 and 777. Every per-seed interval also lies below zero, so the outcome does not depend on which seed is read.

Window geometry of the fold, measured on its 1,799 annotated tiles: long-baseline span 340 days (median; p10 325), acquisition cadence 25 days, usable fraction 1.000 under the long baseline, 0.000 for the single annotated pair and 0.977 for the widened window, whose median span is 30 days. The interval construction therefore behaves as on the other folds; what differs is the target scale and the spectral composition of the scenes, which is why the outcome is reported as a failed replication of the size of the gain rather than as a failure of the window itself.

S1-L. Per-event terrain statistics

The main text reports the elevation and slope ranges of the eleven triggering events in Section 2.1. Both quantities come from the topography stack of the released dataset, which carries five channels: elevation in metres, slope in degrees, aspect, curvature and a topographic wetness index. The arrays are used in their released units, without further normalisation; elevation is bounded to 0-8800 m and slope is computed on the 10 m grid.

For each event, every annotated tile was read and two summaries were taken: the second and 98th percentiles of elevation inside the tile, and the median of those two quantities across the tiles of the event. Aggregating the within-tile percentile rather than the raw pixels keeps a single large tile from dominating an inventory, and the median across tiles is reported instead of the mean for the same reason. Table S1-L gives the result for the 4,392 annotated tiles; the counts per event reproduce the tile counts of Table 1 of the main text, which is a consistency check on the join between the two tables.

Table S1-L. Per-event terrain statistics from the released elevation channel. Tiles are annotated tiles; the elevation columns are the second percentile, the median and the 98th percentile of the within-tile elevation distribution, aggregated over the tiles of the event; slope is the median within-tile slope, aggregated the same way. The eleven events sum to 4,392 tiles, the number reported in Table 1 of the main text.

Event	Tiles	Elev. p2 (m)	Elev. median (m)	Elev. p98 (m)	Slope (deg)
italy	1799	8	325	1260	16.7
chimanimani	345	428	1381	2142	19.3
hiroshima	322	0	274	721	19.3
hokkaido	144	24	162	362	20.8
itogon	82	258	1039	1776	33.0
china	47	918	1947	3166	40.5
indonesia	253	162	1033	2349	31.0

Event	Tiles	Elev. p2 (m)	Elev. median (m)	Elev. p98 (m)	Slope (deg)
kyrgyzstan1	694	846	1674	2907	19.3
kyrgyzstan2	170	1166	2098	2956	24.2
newzealand	508	0	544	1843	31.1
usa (Puerto Rico)	28	0	38	282	7.2
Total	4392	-	-	-	-

Three properties of Table S1-L are used in the main text. First, the tenfold range of event medians (38 m in Puerto Rico, 2,098 m in the second Kyrgyz inventory) and the within-tile range (up to 3,166 m at the 98th percentile in the Chinese inventory) show that the events are not drawn from a single orographic setting. Second, the median slope of a tile spans 7.2 to 40.5 degrees, with nine of the eleven events between 16.7 and 33.0 degrees; the two events outside that band are the flat coastal inventory and the steepest mountain inventory, which is the expected behaviour of a set that spans both. Third, the terrain channels are model inputs, so these are the only conditioning variables that this study both consumes and characterises quantitatively; meteorological variables are neither consumed nor claimed (Section 5.3 of the main text).

Provenance and reproducibility. The statistics are computed by a script archived with the code that reads the topography array of every annotated tile and writes a JSON summary with one row per event. The script and the JSON are part of the data-availability package; re-running it reproduces Table S1-L exactly, because the computation is deterministic and reads the released arrays without any learned parameter.

S1-M. Verification of the event dates used by the rule

The Event-Bracketing Rule takes exactly one input: the event date. Everything else in the rule is read from the acquisition series. The exactness figure reported in the main text (97.93% over 3,135 tiles) is therefore conditional on that date, and it is scored against the annotated pre/post pair of the same dataset that supplies the acquisition dates. Two distinct questions follow, and they are kept separate throughout.

- *Is the input date right?* This can be answered against records outside the released dataset, and it has been answered for every single-event inventory (Table S1-M).
- *Would an independent interpreter have chosen the same pair?* This would require an independently produced inventory for these events. None exists, and the main text says so rather than implying otherwise (Section 5.3).

Table S1-M. Event dates and their verification. Date is the value used as the rule's input and is identical to the event date recorded in the released dataset. “Verified against” names the class of record used for the check; the two earthquake-triggered inventories were checked against the United States Geological Survey event catalogue, where the local date is the one that governs the bracket (for Hokkaido the catalogue time is 18:07 UTC on 5 September 2018, which is 6 September in Japan). The five multi-event inventories are not listed because they carry a date per image rather than per inventory and therefore have no single date to verify.

Event	Date	Verified against	Agreement
Emilia-Romagna (rainfall, flood)	2023-05-16	national civil-protection record	yes
Chimanimani (Cyclone Idai)	2019-03-15	national meteorological record	yes
SW Honshu (heavy rainfall)	2018-06-28	national meteorological record	yes
Hokkaido (earthquake)	2018-09-06	USGS event catalogue (JST date)	yes
Itogon (Typhoon Mangkhut)	2018-09-15	national meteorological record	yes
Puerto Rico (earthquake sequence)	2020-01-07	USGS event catalogue	yes

One event needs a note. The south-western Honshu inventory records 2018-06-28, and its annotated pair runs from 26 June to 1 July. The heavy rain that struck the same region on 5-8 July 2018 is a later episode, and the inventory should not be read as covering it. The date in Table S1-M is the inventory's own, and it is that date that the rule reproduces.

Two settings that must be fixed before comparing. Two conventions have to be fixed before any such check is comparable, and both are easy to get wrong. First, the date that governs the bracket is the local date of the triggering event, not the UTC date of the catalogue entry: the Hokkaido event is 2018-09-05 in UTC and 2018-09-06 in Japan, and the dataset records 2018-09-06. Second, the magnitude of an event is reported on different scales by different agencies; the Hokkaido earthquake is Mj 6.7 on the Japan Meteorological Agency scale and Mww 6.6 in the USGS catalogue, so the scale must be named wherever a magnitude is quoted.

Re-running the check against an independent date. The verification is scripted and takes a table with one row per event: the event name, the date recorded in the dataset, the independent date, the source of that date and the retrieval date. Rows with an empty independent date are skipped and counted as uncovered; nothing is imputed. The script rejoins the table to (i) the per-tile acquisition series and (ii) the annotated pre/post indices, applies the rule (pre = last acquisition strictly before the date, post = first strictly after), and prints exactness and within-one agreement both with the dataset date and with the independent date, per event and overall. Run with the dataset dates it reproduces the published figure (0.9793 exact, 1.0000 within one); run with independent dates it produces the same two numbers for the covered events. The script, the input table and the JSON output are archived with the code.

What this does and does not establish. It establishes that the rule is not being fed a wrong date, which is the failure mode that would invalidate every downstream number while leaving them internally consistent. It does not establish that the annotated window is the only defensible window for these events, and it cannot be extended to the five multi-event inventories, whose annotated windows are tied to per-image dates. Both limits are stated in Section 5.3 of the main text rather than left for the reader to discover.